

Multifidelity Paths and Trees

Dr TP Prescott*

May 5, 2022

I Preliminaries

I.1 Bayesian Inference

Let $f(\cdot \mid \theta)$ denote a high-fidelity model of the data-generating process that has produced data, y_{obs} . The likelihood of the parameters under this model is denoted by $\mathcal{L}(\theta) = f(y_{\text{obs}} \mid \theta)$. For Bayesian inference, we also specify a prior distribution, $\pi(\cdot)$, on the parameter values. The likelihood allows us to update the prior in the context of data into the *posterior*,

$$\mathcal{P}(\theta \mid y_{\text{obs}}) = \frac{\mathcal{L}(\theta)\pi(\theta)}{\mathcal{Z}},$$

where $\mathcal{Z} = \int \mathcal{L}(\theta)\pi(\theta) \, d\theta$ is a normalisation constant, ensuring that the posterior, $\mathcal{P}(\theta \mid y_{\text{obs}})$, is a probability distribution. Finally, given an arbitrary integrable function $G : \theta \mapsto \mathbb{R}$ defined on the parameter space, we denote the *posterior expectation* of G by

$$\mathbb{E}(G \mid y_{\text{obs}}) = \int G(\theta)\mathcal{P}(\theta \mid y_{\text{obs}}) \, d\theta.$$

However, we assume that the likelihood function, $\mathcal{L}(\theta)$, cannot be computed, and therefore that the derived quantities are intractable.

*The Alan Turing Institute, London NW1, England.

1.2 Likelihood-Free Bayes

Given θ , the likelihood-free setting replaces calculation of $\mathcal{L}(\theta)$ with a stochastic estimate, $\omega(y)$, which we will term a *likelihood-free weighting*. This is a random variable, defined on the probability space of simulations of the model, $y \sim f(\cdot \mid \theta)$, under the given parameter. The *simulation-based approximate likelihood* is defined as the conditional expectation,

$$L(\theta) = \mathbf{E}(\omega \mid \theta) = \int \omega(y) f(y \mid \theta) \, dy,$$

of the stochastic likelihood-free weighting. The simulation-based approximate likelihood induces a *simulation-based approximate posterior*,

$$P(\theta \mid y_{\text{obs}}) = \frac{L(\theta)\pi(\theta)}{Z},$$

with normalisation constant $Z = \int L(\theta)\pi(\theta) \, d\theta$. The distribution P induces a *simulation-based approximate posterior expectation* of the arbitrary function G , denoted

$$\mathbf{E}_P(G \mid y_{\text{obs}}) = \int G(\theta)P(\theta) \, d\theta.$$

The *fidelity* of likelihood-free inference is defined by how well $\mathbf{E}_P(G \mid y_{\text{obs}})$ approximates the true posterior, $\mathbf{E}(G \mid y_{\text{obs}})$, and relies on the function $L(\theta)$ being approximately proportional to $\mathcal{L}(\theta)$.

1.2.1 ABC

For example, the approximate Bayesian computation version of the likelihood-free weighting is defined as

$$\omega_{\text{ABC}}(y) = \mathbf{1}(y \in \Omega(y_{\text{obs}}))$$

for an appropriately defined neighbourhood of the observed data, $\Omega(y_{\text{obs}})$. The corresponding simulation-based approximate likelihood is

$$L_{\text{ABC}}(\theta) = \mathbf{P}(y \in \Omega(y_{\text{obs}}) \mid \theta),$$

which is approximately proportional to $\mathcal{L}(\theta)$ for an appropriately chosen neighbourhood, Ω . The corresponding simulation-based approximate posterior induced by L_{ABC} is

$$P_{\text{ABC}}(\theta \mid y_{\text{obs}}) = \frac{L_{\text{ABC}}(\theta)\pi(\theta)}{Z_{\text{ABC}}},$$

for the normalisation constant $Z_{\text{ABC}} = \int L_{\text{ABC}}(\theta)\pi(\theta) \, d\theta$.

1.3 Likelihood-Free Monte Carlo

Consider a weighted Monte Carlo sample produced by Algorithm 1.

Algorithm 1 Stop condition can be computational budget exceeded, or N sufficiently large, etc.

```

Initialise  $N = 0$ .
repeat
  Increment  $N \leftarrow N + 1$ .
  Generate  $\theta \sim \pi(\cdot)$ .
  Simulate  $y \sim f(\cdot \mid \theta)$ .
  Store weight  $w_N \leftarrow \omega(y)$ .
  Store  $\theta_N = \theta$ .
until stop condition met.
return  $\{(w_n, \theta_n)\}_{n=1}^N$ .
```

For any integrable function $G : \theta \mapsto \mathbb{R}$, the ratio

$$\hat{G}_N = \frac{\sum_{n=1}^N w_n G(\theta_n)}{\sum_{m=1}^N w_m} \approx \mathbf{E}_P(G \mid y_{\text{obs}}) \quad (\text{I})$$

approximates the simulation-based approximate posterior expectation of G . The estimate, \hat{G}_N , is consistent as $N \rightarrow \infty$, in the sense that $\hat{G}_N \rightarrow \mathbf{E}_P(G \mid y_{\text{obs}})$ approaches the approximate posterior expectation. In general, there remains bias versus the true posterior expectation, $\mathbf{E}(G \mid y_{\text{obs}})$, as $N \rightarrow \infty$.

2 Multifidelity Inference

Likelihood-free inference as described above relies on repeated simulation of the model of interest, once per iteration in Algorithm 1. If the *cost* of generating $y \sim$

$f(\cdot \mid \theta)$ is prohibitively expensive, this should be required as little as possible.¹ In this section, we allow for alternative, approximate models to be simulated in place of $f(\cdot \mid \theta)$. The key assumption here is that these *low-fidelity* models are significantly less expensive to simulate, at the cost of accuracy. Below, we show how to incorporate such models into likelihood-free inference to produce weighted Monte Carlo samples. As with the weighted Monte Carlo samples produced by Algorithm 1, these samples will be used to estimate approximate posterior expectations, $E_P(G \mid y_{\text{obs}})$, for the simulation-based approximate posterior $P(\theta \mid y_{\text{obs}})$ induced by the high-fidelity model, $f(\cdot \mid \theta)$, and its associated likelihood-free weighting $\omega(y)$.

2.1 Multifidelity Sequential Simulation

Suppose that there exists an ordered list, $f_{\text{lo}} = (f^{(1)}, \dots, f^{(K)})$ of K *low-fidelity* approximations of the *high-fidelity* model, $f^{(K+1)} = f$. We assume for simplicity that each low fidelity model has the same parameter space as the high-fidelity model, and that simulations $y^{(k)} \sim f^{(k)}(\cdot \mid \theta)$ of each low-fidelity model take values in the same output space as simulations $y \sim f(\cdot \mid \theta)$. This assumption can be relaxed with little consequence beyond notational faff.

In the original setting, each iteration of Algorithm 1 depends on a draw from the probability space of parameter–simulation pairs, (θ, y) , with corresponding density $\pi(\theta)f(y \mid \theta)$. Using the additional K low-fidelity models in f_{lo} , we can extend this probability space to parameter–simulation tuples $z^{(K+1)} = (\theta, y^{(1)}, \dots, y^{(K)}, y^{(K+1)})$. Assuming that the simulations in $z^{(K+1)}$ are produced in sequence, we also denote the partially completed sequence by $z^{(k)} = (\theta, y^{(1)}, \dots, y^{(k)})$, for $k = 1, \dots, K$, writing $z^{(0)} = (\theta)$ where no simulations have been completed. The marginal density of each $z^{(k)}$ is given by the recursive relationship

$$g^{(k)}(z^{(k)}) = \begin{cases} f^{(k)}(y^{(k)} \mid z^{(k-1)})g^{(k-1)}(z^{(k-1)}) & k = 1, \dots, K+1, \\ \pi(\theta) & k = 0. \end{cases}$$

Note that we have allowed that simulations $y^{(k)} \sim f^{(k)}(\cdot \mid z^{(k-1)})$ of the k^{th} model, for $k = 2, \dots, K+1$, may depend not only on θ but also all preceding simulations

¹Here, we will take “cost” to be elapsed time, but there may be other useful measurements of cost.

in the sequence, $y^{(1)}$ to $y^{(k-1)}$. This is known as *coupling*; it is not necessary but is useful enough to justify the notational faff.

2.2 Multifidelity Trees

Algorithm 2 defines the function `MF_Tree`, which produces a random *multifidelity tree*, denoted $\tau^{(k)}$ for $k = 0, \dots, K+1$, where $\tau^{(0)} = \text{MF_Tree}()$ and $\tau^{(k)} = \text{MF_Tree}(z^{(k-1)})$ for $k > 0$. Slightly abusing notation, we denote the random output of `MF_Tree` by $T^{(0)} \sim \text{MF_Tree}(\cdot)$ and $T^{(k)} \sim \text{MF_Tree}(\cdot \mid z^{(k-1)})$.

In addition to the prior, π , and the models $f^{(k)}$ for $k = 1, \dots, K+1$, Algorithm 2 requires a positive-valued function, μ , defined on the space of (partial) multifidelity simulations, $z^{(k)}$. This function determines the Poisson distribution of M , the random number of child subtrees of $\tau^{(k)}$, conditional on $z^{(k)}$.

2.3 Multifidelity Likelihood-Free Weighting

Having assumed for simplicity that simulations $y^{(k)} \sim f^{(k)}(\cdot \mid z^{(k-1)})$ of each model take values in the same output space, we assume that the likelihood-free weighting $\omega(y)$, originally defined on the high-fidelity model output, can also be applied to the outputs $y^{(k)}$ from the low-fidelity models.² Then, for any multifidelity tree of the form

$$\tau^{(k)} = \left(z^{(k)}, c^{(k)}, \mu^{(k)}, (\tau_m^{(k+1)})_{m=1}^M \right),$$

we define the recursive function

$$w(\tau^{(k)}) = \begin{cases} \frac{1}{\mu^{(0)}} \sum_{m=1}^M w(\tau_m^{(1)}), & k = 0, \\ \omega(y^{(k)}) + \frac{1}{\mu^{(k)}} \sum_{m=1}^M \left(w(\tau_m^{(k+1)}) - \omega(y^{(k)}) \right), & 1 \leq k \leq K, \\ \omega(y^{(K+1)}), & k = K+1. \end{cases} \quad (2)$$

Note that, by definition, trees $\tau^{(K+1)} = \text{MF_Tree}(z^{(K)})$ have no child subtrees. For any observed multifidelity tree $\tau^{(0)} = \text{MF_Tree}()$, the value of $w(\tau^{(0)})$ is referred to as its *multifidelity likelihood-free weighting*.

²In principle, each low-fidelity model may have its own likelihood-free weighting function $\omega^{(k)}(y^{(k)})$. We neglect this extension for notational simplicity.

Algorithm 2 Recursive generation of a *multifidelity tree* is defined by the overloaded function MF_Tree. Each tree $\tau^{(k)}$ is associated with a realisation of $z^{(k)} = (\theta, y^{(1)}, \dots, y^{(k)})$, produced conditional on $z^{(k-1)}$. The function MF_Tree assigns to this node a Poisson number, M , of child subtrees $\tau_m^{(k+1)}$, where M has conditional mean $\mu(z^{(k)})$. Nodes at depth $K + 1$ are realisations of the full multifidelity simulation, $z^{(K+1)} = (\theta, y^{(1)}, \dots, y^{(K)}, y^{(K+1)})$, and have no child subtrees.

```

function MF_Tree()
    Generate  $\theta \sim \pi(\cdot)$  with observed cost  $c^{(0)}$ .
    Denote  $z^{(0)} = (\theta)$ .
    Set  $\mu^{(0)} = \mu(z^{(0)})$ .
    Generate  $M \sim \text{Poisson}(\mu^{(0)})$ .
    for  $m = 1, \dots, M$  do
        Generate  $\tau_m^{(1)} = \text{MF\_Tree}(z^{(0)})$ 
    end for
    return  $\tau^{(0)} = (z^{(0)}, c^{(0)}, \mu^{(0)}, (\tau_m^{(1)})_{m=1}^M)$ .
end function

```

```

function MF_Tree( $z^{(k-1)}$ ) for  $k = 1, \dots, K$ 
    Simulate  $y^{(k)} \sim f^{(k)}(\cdot \mid z^{(k-1)})$  with observed cost  $c^{(k)}$ .
    Append  $y^{(k)}$  to  $z^{(k-1)}$  to give  $z^{(k)}$ .
    Set  $\mu^{(k)} = \mu(z^{(k)})$ .
    Generate  $M \sim \text{Poisson}(\mu^{(k)})$ .
    for  $m = 1, \dots, M$  do
        Generate  $\tau_m^{(k+1)} = \text{MF\_Tree}(z^{(k)})$ .
    end for
    return  $\tau^{(k)} = (z^{(k)}, c^{(k)}, \mu^{(k)}, (\tau_m^{(k+1)})_{m=1}^M)$ .
end function

```

```

function MF_Tree( $z^{(K)}$ )
    Simulate  $y^{(K+1)} \sim f^{(K+1)}(\cdot \mid z^{(K)})$  with observed cost  $c^{(K+1)}$ .
    Append  $y^{(K+1)}$  to  $z^{(K)}$  to give  $z^{(K+1)}$ .
    Return  $\tau^{(K+1)} = (z^{(K+1)}, c^{(K+1)}, 0, ( ))$ .
end function

```

RESULT The random multifidelity tree $T^{(0)} \sim \text{MF_Tree}()$ induces a random multifidelity likelihood-free weight, $W = w(T^{(0)})$. Given θ , the conditional expectation of W is

$$\mathbf{E}(W \mid \theta) = \mathbf{E}(\omega(y) \mid \theta) = L(\theta),$$

equal to the high-fidelity simulation-based approximation to the likelihood.

COROLLARY For the weighted Monte Carlo sample produced by Algorithm 3, the ratio estimate \tilde{G}_N defined in Equation (1) is a consistent estimate of the high-fidelity simulation-based approximate posterior expectation $\mathbf{E}_P(G \mid y_{\text{obs}})$.

Algorithm 3 Stop condition can be computational budget exceeded, or N sufficiently large, etc.

Initialise $N = 0$.

repeat

Increment $N \leftarrow N + 1$.

Generate $\tau^{(0)} = \text{MF_Tree}()$ of form $\tau^{(0)} = (z_0, c_0, \mu_0, (\tau_m^{(1)})_{m=1}^M)$.

Store weight $w_N \leftarrow w(\tau^{(0)})$ for w in Equation (2).

Store $\theta_N \leftarrow \theta$ from $z_0 = (\theta)$.

until stop condition met.

return $\{(w_n, \theta_n)\}_{n=1}^N$.

2.4 Performance

The results above are true for any function, μ , defining the conditional Poisson mean of M in a call of $\text{MF_Tree}()$ or $\text{MF_Tree}(z^{(k-1)})$, so long as μ takes strictly positive values. We assume that the computational cost of an iteration of Algorithm 3 is dominated by generating $T^{(0)} \sim \text{MF_Tree}(\cdot)$. The recursive function

$$C(\tau^{(k)}) = c^{(k)} + \sum_{m=1}^M C(\tau_m^{(k+1)}),$$

represents the observed cost of generating $\tau^{(k)} = \text{MF_Tree}(z^{(k-1)})$, where we note that $C(\tau^{(K+1)}) = c^{(K+1)}$. It follows that the cost of a given iteration of Algorithm 3 is the cost, $C(\tau^{(0)})$, of the corresponding multifidelity tree, $\tau^{(0)} = \text{MF_Tree}()$. The distribution of $C(T^{(0)})$ for $T^{(0)} \sim \text{MF_Tree}(\cdot)$ thus depends on the function

μ through its influence on M .

Smaller values of μ reduce the computational cost of generating $T^{(0)} \sim \text{MF_Tree}(\cdot)$, by reducing the number of calls of $\text{MF_Tree}(z_m^{(k)})$ within the inner loop of any call of $\text{MF_Tree}(z^{(k-1)})$. However, small values for μ incur less high quality Monte Carlo samples, in the sense of increasing the mean squared error (MSE) of \tilde{G}_N as an estimate of $\tilde{G} = \mathbf{E}_P(G \mid y_{\text{obs}})$. We can derive the leading order behaviour of the MSE as a function of the total computational cost, C_{tot} , of producing a sample from Algorithm 3 as

$$\text{MSE} = \frac{\mathbf{E}((\Delta(T)w(T))^2)\mathbf{E}(C(T))}{\mathbf{E}(w(T))^2} \frac{1}{C_{\text{tot}}} + O(C_{\text{tot}}^{-2})$$

for the random multifidelity tree $T \sim \text{MF_Tree}(\cdot)$. Here w and C are given above, and $\Delta(\tau) = G(\theta) - \tilde{G}$ for the approximate posterior expectation $\tilde{G} = \mathbf{E}_P(G \mid y_{\text{obs}})$. We write

$$\mathcal{J}[\mu] = \mathbf{E}((\Delta(T)w(T))^2)\mathbf{E}(C(T))$$

as a functional of μ , and note that $\mathbf{E}(w(T)) = Z$ is equal to the normalisation constant, independently of μ . Thus, minimising $\mathcal{J}[\mu]$ will optimise the performance of Algorithm 3, in the sense of producing a Monte Carlo sample with the smallest MSE for a given computational budget.³

The functional $\mathcal{J}[\mu]$ is, in general, difficult to work with. In the following section we will define a data-driven cost functional $J[\mu]$ to work with instead.

3 Multifidelity Implementation

3.1 Cost functionals

3.1.1 Simulation cost functional

Consider any one multifidelity tree

$$\tau^{(0)} = \left(z^{(0)}, c^{(0)}, \mu^{(0)}, \left(\tau_m^{(1)} \right)_{m=1}^M \right) = \text{MF_Tree}()$$

³Equivalently, producing a Monte Carlo sample with a given MSE as cheaply as possible.

produced for an iteration of Algorithm 3. We assume that a baseline mean function, μ , is used for the call of `MF_Tree()`. For any proposed alternative mean function, μ' , we can define the recursive function

$$\begin{aligned} C(\tau^{(k)}; \mu') &= c^{(k)} + \frac{\mu'(z^{(k)})}{\mu^{(k)}} \sum_{m=1}^M C(\tau_m^{(k+1)}; \mu'), \\ C(\tau^{(K+1)}; \mu') &= c^{(K+1)}. \end{aligned}$$

The function $C(\tau^{(k)}; \mu')$ is a data-driven estimate of the expected cost of producing $\tau^{(k)} = \text{MF_Tree}(z^{(k-1)})$, using μ' as the mean function in place of μ .

Now consider the the set of multifidelity trees $\tau_n^{(0)}$, where $n = 1, \dots, N$, produced by N independent calls of `MF_Tree()` in Algorithm 3. We define the data-driven simulation cost functional⁴

$$C[\mu'] = \frac{1}{N} \sum_{n=1}^N C(\tau_n; \mu')$$

as the average value of $C(\tau_n; \mu')$ for the sample of independent observations, $\tau_n^{(0)}$, of the random multifidelity tree, $T^{(0)} \sim \text{MF_Tree}(\cdot)$.

3.1.2 MSE cost functional

Under the same assumptions as above, we define the recursive function

$$\begin{aligned} V(\tau^{(k)}; \mu') &= \left(\frac{1}{\mu^{(k)}} \right)^2 \left(\sum_{m_1 \neq m_2} \left(w(\tau_{m_1}^{(k+1)}) - \omega(y^{(k-1)}) \right) \left(w(\tau_{m_2}^{(k+1)}) - \omega(y^{(k-1)}) \right) \right. \\ &\quad \left. + \frac{\mu^{(k)}}{\mu'(z^{(k)})} \sum_{m=1}^M V(\tau_m^{(k+1)}; \mu') \right), \end{aligned}$$

⁴ C is doing a lot of work here!

valid for $k = 2, 3, \dots, K$, with the ‘boundary conditions’,

$$\begin{aligned} V(\tau^{(0)}; \mu') &= \left(\frac{1}{\mu^{(0)}} \right)^2 \left(\sum_{m_1 \neq m_2} w(\tau_{m_1}^{(1)}) w(\tau_{m_2}^{(1)}) + \frac{\mu^{(0)}}{\mu'(z^{(0)})} \sum_{m=1}^M V(\tau_m^{(1)}; \mu'), \right) \\ V(\tau^{(1)}; \mu') &= \left(\frac{1}{\mu^{(1)}} \right)^2 \left(\sum_{m_1 \neq m_2} w(\tau_{m_1}^{(2)}) w(\tau_{m_2}^{(2)}) + \frac{\mu^{(1)}}{\mu'(z^{(1)})} \sum_{m=1}^M V(\tau_m^{(2)}; \mu'), \right) \\ V(\tau^{(K+1)}; \mu') &= \left(\omega(y^{(K+1)}) - \omega(y^{(K)}) \right)^2. \end{aligned}$$

Note that if a given tree $\tau^{(k)}$ has no child subtrees, then $V(\tau^{(k)}; \mu') = 0$.

Now consider the set of multifidelity trees $\tau_n^{(0)}$, where $n = 1, \dots, N$, produced by N independent calls of `MF_Tree()` in Algorithm 3. We define the data-driven MSE cost functional

$$V[\mu'] = \frac{1}{N} \sum_{n=1}^N \left(G(\theta_n) - \bar{G}_N \right)^2 V(\tau_n^{(0)}; \mu'),$$

where θ_n is the parameter associated with $\tau_n^{(0)}$ and where \bar{G}_N is the Monte Carlo posterior estimate.

3.1.3 Performance cost functional

In summary, if we take the multifidelity trees produced by a batch run of Algorithm 3, then we can construct a cost functional for a new mean function μ' , equal to

$$J[\mu'] = C[\mu'] V[\mu'].$$

This cost functional can now be used to update the mean function. In the following, we will construct a mean function as a recurrent neural network, and train it on simulation data, based on minimising this cost function.

3.2 Recurrent Neural Network μ

Layers X , F and G define the dynamic system:

- $x^{(0)} = X(\theta)$
- $x^{(k)} = F(x^{(k-1)}, y^{(k)})$ for $k = 1, \dots, K$.

- $\mu(z^{(k)}) = G(x^{(k)})$, where we use a positive activation function such as `softplus` to ensure $\mu > 0$.

The details (i.e. sizes) of these layers will be application specific.

We train these layers by minimising a regularised cost function

$$J[\mu] + \gamma R[\mu]$$

derived from an existing batch of multifidelity simulations produced by Algorithm 3. Here, $J[\mu]$ is the performance cost functional derived in the previous subsection, which trades MSE against simulation cost. The regulariser functional, $R[\mu]$, needs to be chosen to ensure that μ does not take excessively small values. For example, we can use L^2 weight regularisation of the output layer $G : x^{(k)} \mapsto \mu(z^{(k)})$. If for a given positive β the layer G uses the `softplus` activation function,

$$\sigma_{\beta}(t) = \frac{1}{\beta} \log(1 + \exp(\beta t)) ,$$

then this regularisation can ensure a prior on μ to take values the order of $(\log 2)/\beta$. Note that the strength of the regularisation is parametrised by the free hyperparameter $\gamma \geq 0$.